Prompt■Adaptation Methods: Cost & VRAM Cheat■Sheet

Method	Trainable Params	VRAM (T4 16GB)	1■Epoch Time (SST■2)
Soft Prompt (20 tokens)	≈15k	~1.5■GB	~1■min
Prefix■Tuning (100 tokens)	≈75k	~2■GB	~3■min
LoRA r=8 (GPT■2)	≈4.5■M	~4■GB	~8■min
Full Fine■Tune GPT■2	124■M	~8■GB	~25■min

**Approx. wall■clock on free Colab T4, batch=4, fp16 when possible.*

Use as rule■of■thumb to pick the quickest viable technique.